

The Phoenix transaction model

Dusk

Last update: November 15, 2023

Abstract. Recent advances in blockchain technologies have led to their increased usage in different use cases, with one prominent case being cryptocurrencies. Often, this involves large amounts of transaction data being shared over the network. Even when the public keys of the users are not directly associated with identities, as is the case in Bitcoin, sending this data in plaintext can lead to traceability and de-anonymization. To solve this problem, projects like Zcash introduced a private-by-design blockchain capable of issuing obfuscated (private) transactions, where the receiver, the sender, and the amount of exchanged money remain private. This is possible thanks to zero-knowledge proofs, with the drawback of the high computing resources they require. This document gives an overview and discussion of Phoenix, the transaction model used by Dusk. Phoenix is inspired by the Zcash transaction model, sharing many of its privacy features, and includes a novel delegation model in which users can delegate proof computations to partially trusted third parties.

1 Introduction

One classical problem in cryptocurrencies is anonymity. Since the deployment of the Bitcoin blockchain in 2009 [26], many concerns about the security and privacy guarantees of such networks have arisen. In most transparent blockchains, transaction information like the addresses of sender and receiver, and the amount of money being transferred, are public. Although an address is not directly linked to a user's identity, the transparency of the whole system makes transactions traceable. Furthermore, questions about Bitcoin's privacy have spanned an increasing body of research, and many studies show that it is feasible to identify and track participants of the Bitcoin network [30,18]. Whereas there have been attempts at mitigating the traceability of coins through the use of multiple addresses and mixers [4], this approach has only found limited success.

Security of transactions in this context is at higher risk than with traditional banking: if an individual gets ever linked to its address, privacy of all past and future movements associated with that account is totally lost. Furthermore, the interest of society and regulators in the topic has grown in the recent years, both for individual use and industry trade secrets. The need to guarantee the right to privacy has ended up with new legislation like the General Data Protection Regulation (GDPR) of the European Union [1] or the California Consumer Privacy Act (CCPA) [2].

However, it is a considerable challenge to use a transparent system like the Bitcoin blockchain and at the same time provide privacy to sensitive data. Recently, a great effort has been put in developing alternative models with anonymity baked into the design, the most prominent examples of this approach being Zcash [32,19] and Monero [34]. These systems allow for conducting private transactions over decentralized networks. Zcash built a solution based on the *unspent transaction output (UTXO)* model used in Bitcoin, but achieving anonymity at the same time. Zcash has an on-and-off privacy switch, whereas transactions in Monero are always private.

While initial cryptocurrencies used relatively mild cryptographic primitives, namely hashes and signature schemes, these newer private coins leverage heavy machinery, including very efficient zero-knowledge proof systems (zk-SNARKs) [17,5,14] or ring signatures [15,27]. This allows for transactions (called *shielded transactions* in the context of Zcash) in which sender address, receiver address and value are all hidden from the public, and yet the transactions can still be registered in the blockchain.

One recent project is Dusk [24], an open-source public blockchain network designed for securities trading and other financial applications. In this context, privacy is not only important, but it is essential for businesses and regulated markets. From a technical perspective, the main difference with other privacy-oriented blockchains is that Dusk aims to support the capability of executing confidential smart contracts [23].

Contributions. This document gives an overview of *Phoenix* [9], the transaction model developed by Dusk, capable of issuing obfuscated (private) transactions across Dusk’s blockchain. It is a UTXO model based on Zcash [19] and CryptoNote [34], with the addition of a method for delegating the expensive proof computations to a partially trusted helper. The goal of this document is to explain the different aspects of the protocol, from the concrete cryptographic building blocks and their parameter choices, to the structure and flow of the transactions, giving an intuition on security and the reasoning behind the decisions made.

Roadmap. The paper is organized as follows. In Section 2, we give a high-level overview of the transactions flow. In Section 3, we present the set of cryptographic primitives used by Phoenix. In Section 4, we describe Phoenix within the Dusk context. We first describe the network, and then how transactions are created and processed. In Section 5, we analyze the protocol and its properties, giving an intuition on security, and its implementation, for which we provide benchmarks.

2 Model intuition

Before we introduce the Phoenix protocol in detail, we give a simplified example to illustrate the ideas behind it. Suppose there are three parties Alice, Bob, and

Charlie, that want to exchange some money secretly. In particular, Alice wants to send some money v to Bob, and then Bob wants to transfer the same amount of money to Charlie using obfuscated transactions.

$$\text{Alice} \xrightarrow{v} \text{Bob} \xrightarrow{v} \text{Charlie}$$

We show how the protocol works in two steps. First, we explain how Alice would send v to Bob, without explaining how she got the money in the first place, and afterwards, we focus on how Bob can spend the money received from Alice to send it to Charlie.

Alice \rightarrow Bob

If Alice wants to send v to Bob, Alice must create a new note N_B of value v that only Bob can spend, and sent it to the network as part of a transaction. The note N_B will contain the following fields:

$$N_B = \{\text{value commitment, encrypted data, note public key}\}.$$

The value commitment is a commitment to the amount of money v that Alice transfers to Bob. The encrypted data is encrypted using Bob's public key, so that Bob can decrypt it to recover the opening of the value commitment. The note public key is a one-time public key that will allow Bob to spend the note; it is derived from Bob's public key in a way that only Bob (not even Alice), can compute its corresponding secret key.

Once Alice's transaction that includes N_B is accepted by the network, the note N_B is included as a new leaf in the Merkle tree of notes, a data structure that allows for efficient membership proofs.

Bob \rightarrow Charlie

Assume that now Bob wants to spend the money v from N_B . Bob must show that N_B is a valid note that is included in the Merkle tree of notes, and ownership of the note (i.e. knowledge of the private key associated with the note public key). Since we want the transaction to be private, we don't want to disclose which note nor how much money Bob is spending. Thus, we use a zero-knowledge proof to ensure that:

1. The note N_B is stored in the Merkle tree of notes (membership).
2. Bob knows the secret key associated with the note public key of N_B (ownership).

To prevent Bob from spending N_B more than once, the system requires from Bob a value, called nullifier, as part of the transaction, which is a unique deterministic value correlated to Bob's secret key and the note, but cannot be linked to the note by external parties (i.e. still hides which specific note Bob is spending). If Bob tried to spend N_B again, they would get the same nullifier, and

since it must be included as part of the transaction, the network would detect the attempt to spend a note that is already nullified. This means that the network must keep track of all nullifiers. Bob has to guarantee that the nullifier has been computed correctly, by providing a zero-knowledge proof of the following statement:

3. The nullifier associated with N_B is computed as it should (nullification).

To send the money from N_B to Charlie, Bob will create a note N_C with value v that only Charlie will be able to spend. The idea is the same as in the previous section, but Bob must also prove that the transfer contains the right amount of money spent, and not more. Therefore, the zero-knowledge proof also proves the following statements:

4. The value commitment associated with N_C is computed as it should (minting).
5. The money sent in N_C is the money spent from N_B (balance integrity).

Thus, Bob sends a transaction to the network including the following fields:

$$\text{Transaction from Bob} = \{\text{spend proof, nullifier, } N_C\}.$$

We remark that, in fact, Alice also included all the above fields in the transaction, but we only focused in the new note's minting part for the sake of clarity.

3 Cryptographic primitives

In this section, we detail the cryptographic primitives used in Phoenix. We briefly introduce Merkle trees, the commitment scheme, encryption scheme, proof system, elliptic curves and hash functions used, specifying at each step the concrete parameters with which each of the primitives is instantiated.

Notation. Throughout the document, we use the following conventions. Given a set S , we denote sampling an element x uniformly at random from S by $x \leftarrow S$. Any group \mathbb{G} used is of a large prime order, and we assume that the discrete logarithm problem is hard in \mathbb{G} . If two elements are denoted by the same letter in upper case and lower case, e.g. a, A , this often signifies the fact that A is a public key corresponding to the secret key a .

3.1 Merkle trees

A *Merkle tree* [25] is a tree that contains at every vertex the hash of its children vertices. More precisely, we consider a perfect k -ary tree of height h . The single vertex at level 0 is called the *root* of the tree, and the k^h vertices at level h are called the *leaves*. Given a vertex in level i , the k vertices in level $i + 1$ that are adjacent to it are called its *children*. Two vertices are each other's *sibling* if they are children of the same vertex.

To each vertex in the tree, we will recursively associate a value, starting from the leaves.¹ Let H be a hash function.

- Level h : leaves are initialized to a null value. Through the lifetime of the tree, they will progressively be filled from left to right with values.
- Level $0 \leq i < h$: each vertex has k children c_1, \dots, c_k at level $i + 1$. We set the value of the vertex to $H(c_1, \dots, c_k)$.

The tree is updated every time a new value is written into a leaf, by updating the $h + 1$ elements in the path from the new value to the root. In particular, this means that the root changes after every update. A nice feature of Merkle trees is that, given a root r , it is easy to prove that a value x is in a leaf of a tree with root r . The proof works as follows:

- *Prove.* For $i = h, \dots, 1$, let x_i be the vertex that is in level i and is in the unique path from x to the root. Let $y_{i,1}, \dots, y_{i,k-1}$ be the $k - 1$ siblings of x_i . Output

$$(x, (y_{1,1}, \dots, y_{1,k-1}), \dots, (y_{h,1}, \dots, y_{h,k-1})).$$
- *Verify.* Parse input as $(x_h, (y_{1,1}, \dots, y_{1,k-1}), \dots, (y_{h,1}, \dots, y_{h,k-1}))$, where x_h is the purported value and $y_{i,1}, \dots, y_{i,k-1}$ are the purported siblings at level i . For $i = h - 1, \dots, 0$, compute²

$$x_i = H(x_{i+1}, y_{i+1,1}, \dots, y_{i+1,k-1}).$$

This allows for proving membership in a set of size k^h by sending $O(kh)$ values. This proof is sound provided that the hash function is collision resistant [21, Section 5.6.2]. For our application, we will set $k = 4$ and $h = 17$.

3.2 Zero-knowledge proof systems

Phoenix uses the zk-SNARK PlonK [14] as its proof system. PlonK allows anyone to prove satisfiability of any arithmetic circuit modulo a prime. Since arithmetic circuit satisfiability is an NP-complete problem, this proof system will allow us to prove any statement in NP. PlonK makes use of the KZG polynomial commitment scheme [20], as described in [14]. This requires instantiating PlonK over a pairing-friendly group, which is described in Section 3.3.

Below is a summary the efficiency of PlonK, for a circuit with n multiplication gates and ℓ public inputs.

- *Proving time:* $O(n)$ group and field operations.
- *Verification time:* $O(1 + \ell)$ group and field operations.
- *Proof size:* $O(1)$ group and field elements.

PlonK is sound in the algebraic group model [13], and statistically zero-knowledge. A complete and explicit description of the scheme can be found in [14, Section 8].

¹ We will often abuse notation and write the vertex to refer to the value associated with the vertex.

² To be precise, the prover also has to send $\lceil \log_2 k \rceil$ bits for each level, specifying the position of x_i with respect to its siblings, so that the verifier knows in which order to arrange the inputs of the hash.

3.3 Elliptic curves

BLS12-381 [6] and Jubjub [7] are the elliptic curves used. More precisely, let

$$q = 4002409555221667393417789825735904156556882819939007885332058136 \\ 124031650490837864442687629129015664037894272559787,$$

$$p = 5243587517512619047944774050818596583769055250052763782260365869 \\ 9938581184513.$$

Note that both are prime numbers, with bit-lengths 381 and 255, respectively. The curve BLS381-12 is the curve over \mathbb{F}_q defined by the equation

$$E : Y^2 = X^3 + 4.$$

We have that $E(\mathbb{F}_q)$ has different subgroups $\mathbb{G}_1, \mathbb{G}_2$ such that $\#\mathbb{G}_1 = \#\mathbb{G}_2 = p$. This curve is pairing-friendly (with embedding degree $k = 12$), so pairings are efficiently computable. More precisely, Phoenix makes use of the bilinear group

$$\mathbb{B} = (p, \mathbb{G}_1, \mathbb{G}_2, \mathbb{G}_T, e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T).$$

By instantiating the zk-SNARK with the bilinear group \mathbb{B} , we are able to prove statements about satisfiability of arithmetic circuits over \mathbb{F}_p , the so-called *scalar field* of E .

Furthermore, we are interested in proving certain operations with the zk-SNARK, like the correct verification of a Schnorr signature σ . Note that σ is an element of a certain elliptic curve J , but is represented as two coordinates in the base field \mathbb{F}_s of J . Therefore, the verification can be best represented as arithmetic constraints modulo s . While it is possible to represent any NP statement using arithmetic modulo p to plug it into the zk-SNARK, this incurs into a significant efficiency loss if not done carefully. The natural thing is to set $s = p$. Therefore, the signature scheme must be instantiated with an elliptic curve over \mathbb{F}_p . For this, let

$$d = -\frac{10240}{10241} \bmod p.$$

Phoenix uses the Jubjub curve, defined by the equation

$$J : -X^2 + Y^2 = 1 + dX^2Y^2,$$

over \mathbb{F}_p . In particular, it uses a subgroup \mathbb{J} of order

$$t = 6554484396890773809930967563523245729705921265872317281365359162 \\ 392183254199,$$

which is a 252-bit prime.

The primes and groups defined here will be used through the rest of the document.

3.4 Hash functions

Phoenix uses two hash functions, depending on the context. Sometimes, we will want to prove a statement involving a hash function H , e.g. given y , prove knowledge of x such that $H(x) = y$. We will do so with PlonK, which requires statements to be written as arithmetic constraints modulo a large prime p . Most hash function evaluations do not naturally translate to this language, incurring in a big efficiency loss.

To avoid this, the Poseidon hash function [16] $H : \mathfrak{F}_p \rightarrow \mathbb{F}_p$, where \mathfrak{F}_p is the set of tuples of \mathbb{F}_p -elements of any length, will be used whenever we compute a hash of which we need to produce a proof. This is because Poseidon is purposefully designed to work with modular arithmetic. A symmetric encryption scheme [22] based on Poseidon is also used, as described below.

Poseidon is built by applying the sponge construction [3] to a permutation $\pi : \mathbb{F}_p^t \rightarrow \mathbb{F}_p^t$, for $t = r + c$, where

- r is the *rate*, i.e. the amount of \mathbb{F}_p -elements of the input that can be processed in a call to π .
- c is the *capacity*, which is a part of the permutation that is never output by the hash, and is required for security.

The permutation π is composed of linear (matrix multiplication over \mathbb{F}_p) and non-linear (S-boxes) operations. Some rounds are *full rounds*, and apply S-boxes to the whole input, and others are *partial rounds*, in which an S-box is applied to a single \mathbb{F}_p -element.

In this case, Phoenix uses POSEIDON-128 to target 128-bit security. Following the recommendations of [16], parameters are set as $r = 4$ and $c = 1$, so that a hash in the Merkle tree can be computed with a single call to the permutation. Internally, a permutation performs $R_F = 8$ full rounds and $R_P = 59$ partial rounds, and uses $S(x) = x^5$ as the S-box.

When no proofs involving hashes are required, the BLAKE2 family of hash functions [31] is used, which yields better efficiency. In particular, the BLAKE2b hash function is used, since it is optimized for 64-bit platforms.

The encryption scheme [22] is a variation of the one-time pad encryption in the field. It uses Poseidon as a pseudorandom function to extend an agreed-upon symmetric key and encrypt the message. Therefore, the encryption scheme is perfectly secure under the random oracle model.

Concretely, the encryption works as follows. Given a message in \mathbb{F}_q^ℓ , each \mathbb{F}_q -component is added to the corresponding component of the key. The key is obtained using Poseidon by extending the symmetric-encryption key, which is an elliptic curve point, to obtain a key with the same size of the message. The initialization vector contains the two coordinates of the key and a nonce, it is passed to the Poseidon iteration which extends the key and outputs the ciphertext. The sender sends the encryption along with the nonce and the information needed to compute the key, so the receiver can use Poseidon with the same key and nonce to decrypt the message.

3.5 Commitments

Phoenix makes use of the well-known Pedersen commitment scheme [28], which we now describe.

- *Setup*. Sample and output the commitment key $\text{ck} = (G, G') \leftarrow \mathbb{J}^2$.
- *Commit*. On input a value v , sample randomness $r \leftarrow \mathbb{F}_t$ and output

$$c = \text{Com}_{\text{ck}}(v; r) = vG + rG'.$$

- *Open*. Reveal v, r . With these values, anyone can recompute the commitment and check whether it matches the commitment previously provided.

The Pedersen commitment scheme is perfectly hiding, and computationally binding under the discrete logarithm assumption.

3.6 Signatures

The Schnorr Sigma protocol [33] is also used, compiled with the Fiat–Shamir transformation [12,29], as a signature scheme. In particular, Phoenix makes use of a double-key version to be able to delegate computations later in the protocol. Let $G, G' \leftarrow \mathbb{J}$.

- *Setup*. Sample a secret key $\text{sk} \leftarrow \mathbb{F}_t$ and set the corresponding public key $(\text{pk}, \text{pk}') = (\text{sk}G, \text{sk}G')$. Output $(\text{sk}, (\text{pk}, \text{pk}'))$.
- *Sign*. On input a message m and a secret key sk , sample $r \leftarrow \mathbb{F}_t$ and compute $(R, R') = (rG, rG')$. Compute the challenge $c = H(m, R, R')$, and set

$$u = r - c\text{sk}.$$

Output the signature $\sigma = (R, R', u)$.

- *Verify*. On input a public key pk , message m and signature $\sigma = (R, R', u)$, compute $c = H(m, R, R')$ and check whether the following equalities hold:

$$\begin{aligned} R &= uG + c\text{pk}, \\ R' &= uG' + c\text{pk}'. \end{aligned}$$

If so, accept the signature, otherwise reject.

The signature scheme is existentially unforgeable under chosen-message attacks under the discrete logarithm assumption, in the random oracle model [21, Section 12.5.1]. While the Schnorr signature scheme is widely known, the double-key version has not been used before, to the best of our knowledge. In the transaction protocol, this is leveraged to allow for delegation of proof computations without the need to share one’s secret key with the helper.

4 Transaction model

Phoenix is the transaction model used by Dusk, an open-source public blockchain with a UTXO-based architecture that allows the execution of obfuscated transactions and confidential smart contracts. In blockchains with a state-account configuration such as Ethereum, a state machine describes the amount of money belonging to each account in the network, and the state changes after validating different sets of transactions. In privacy-preserving blockchains, there are no accounts or wallets at the protocol layer. Instead, coins are stored as a list of UTXOs with a quantity and some criteria for spending it. In this approach, transactions are created by consuming existing UTXOs and producing new ones in their place. Dusk follows this system, and UTXOs are called *notes*. From now on, we also use this terminology.

Unlike transparent transaction models, such as the Bitcoin network, where it is easy to monitor which notes were spent, this task is much harder in a privacy-preserving network, since the details of the notes must be kept hidden. In this case, the network must keep track of all notes ever created by storing their hashes in the leaves of a Merkle tree (called *Merkle tree of notes*). That is, when a transaction is validated, the network includes the hashes of the new notes to the leaves of this tree. To prevent double spending, transactions include a list of deterministic values called *nullifiers*, one for each note being spent, which invalidates these notes. The network nodes must check that nullifiers were not used before. The idea here is that the nullifier is computed in such a way that an external observer cannot link it to any specific note. This way, when a transaction is accepted, the network knows that some notes are nullified and can no longer be spent, but does not know which ones.

At a technical level, transactions consist of two parts, a header which includes metadata, and the payload with the actual contents of the transaction. The sender first sets $\text{tx}_{\text{skeleton}}$, which they use to compute tx_proof . The validator will add the positions of the newly minted notes when the transaction is accepted.

$$\text{tx}_{\text{metadata}} = [\text{timestamp}, \text{block_height}, \text{status}, \text{type}, \text{tx_hash}, \text{tx_fee}, \\ \text{gas_price}, \text{gas_limit}, \text{gas_spent}],$$

$$\text{tx}_{\text{skeleton}} = [\text{root}, \text{nullifiers}: [\text{nullifier}_1, \dots, \text{nullifier}_r], \text{notes}: [\text{N}_1^{\text{new}}, \dots, \text{N}_s^{\text{new}}], \\ \text{fee}: [\text{gas_limit}, \text{gas_price}, \text{gas_change_stealth_address}]].$$

$$\text{tx}_{\text{payload}} = [\text{tx}_{\text{skeleton}}, \text{tx_proof}, \text{positions}: [\text{pos}_1, \dots, \text{pos}_s]].$$

The value tx_hash is the BLAKE2b hash of $\text{tx}_{\text{skeleton}}$. The remaining parameters included in $\text{tx}_{\text{metadata}}$ are self-explanatory. The payload contains the following: **root** is the root of the Merkle tree of notes; **nullifiers** is a list of deterministic values that nullify the notes being spent; **notes** is the list of new notes being minted; **fee** contains the details of the transaction fees, where **gas_change_stealth_address**

is a stealth address (a pair of the form (npk, R)) used to return to the user the difference between gas_limit and gas_spent ; pos is the position of the note in the Merkle tree of notes, and tx_proof is a zero-knowledge proof that proves that the transaction has been performed following the network rules. We provide greater details about this proof in Section 4.2.

All notes in the network have the same structure. More specifically, a note is a data object with the following structure.

$$N = \{\text{type}, \text{com}, \text{nonce}, \text{enc}, \text{npk}, R\}.$$

The parameters above correspond to the following: **type** indicates the type of the note, either transparent or obfuscated; **com** is a commitment to the value of the note; **nonce** is an initialization vector needed for the encryption scheme; **enc** is an encryption of the opening of **com** that can be decrypted using the receiver’s view key; **npk** is the note’s public key, whose associated private key **nsk** can only be computed by the receiver of the note; and R is a point in the Jubjub subgroup \mathbb{J} that allows the receiver to compute **nsk** and also identify that he is the receiver of the transaction. A note will have an associated position, but we do not consider the position part of the note, as the position cannot be assigned by the sender.

4.1 Protocol keys

In this section, we introduce all the different types of keys involved in the protocol. In Phoenix, each note is associated with a unique public key, instead of using the static public key of the receiver. This hinders traceability (the approach was introduced in [34]). Also, it uses two-element keys, which allows users of the network to delegate the process of scanning for the notes addressed to them.

Let $G, G' \leftarrow \mathbb{J}$, as defined in Section 3.3. We denote by H^{Poseidon} and H^{BLAKE2b} the Poseidon and BLAKE2b hash functions, respectively, introduced in Section 3.4. On the one side, every user keeps the following static keys.

- *Secret key*: $\text{sk} = (a, b)$, where $a, b \leftarrow \mathbb{F}_t$.
- *Public key*: $\text{pk} = (A, B)$, where $A = aG$ and $B = bG$.
- *View key*: $\text{vk} = (a, B)$.

On the other side, we assign a one-time key pair to each note issued in the network, computed using the Diffie–Hellman key exchange protocol [8]. The receiver’s Diffie–Hellman partial key will be the first part of its public key, A , whereas the sender will use a fresh key. More precisely, the note public key of a note sent to a receiver with public key pair $\text{pk} = (A, B)$ is computed by the sender as follows:

- *Note public key*: **npk** is computed by the sender as follows.
 1. Sample r uniformly at random from \mathbb{F}_t .
 2. Compute a symmetric Diffie–Hellman key $k_{\text{DH}} = rA$.
 3. Compute a note public key $\text{npk} = H^{\text{BLAKE2b}}(k_{\text{DH}})G + B$.
 4. Compute $R = rG$.

The sender of a note will attach to it the note public key npk and the partial Diffie–Hellman R used to create npk .³ Given a pair (npk, R) , the receiver can identify whether the note was sent to them by recomputing $\tilde{k}_{\text{DH}} = aR$ (using their secret a), and checking the equation

$$\text{npk} \stackrel{?}{=} H^{\text{BLAKE2b}}(\tilde{k}_{\text{DH}})G + B.$$

Note that this check can be performed using only the receiver’s view key, and not their whole secret key. This way, a network user can delegate the job of scanning the different transactions of the network to retrieve their notes by sharing their view key vk with an external entity, which we call a *network-listening helper*. Thus, the user could delegate the scanning of all transactions to a different entity by sharing (a, B) with the helper. Even with that information, such an entity could not spend \mathcal{R} ’s money, since they can not derive $\text{sk}_{\mathcal{R}}$ without the second part of \mathcal{R} ’s private key.

- *Note secret key*: $\text{nsk} = H^{\text{BLAKE2b}}(k_{\text{DH}}) + b$. Note that this key can only be computed by the receiver of the note, since they are the only ones holding the whole secret key $\text{sk} = (a, b)$, and sk cannot be recovered from public information. This is due to the discrete logarithm assumption in \mathbb{J} . The receiver will use the note secret key when spending the note.

When spending a note, the sender will use a variation of the note public key called *nullification key* computed as $\text{npk}' = \text{nsk}G'$.

4.2 Phoenix

We describe Phoenix in the general scenario in which a sender wishes to send different amounts of money v_1, \dots, v_r to different receivers with public keys $\text{pk}_1, \dots, \text{pk}_r$. We assume the sender has enough funds, i.e. they own a set of notes $\{\text{N}_1^{\text{old}}, \dots, \text{N}_s^{\text{old}}\}$ each with an associated amount w_i such that

$$\sum_{i=1}^s w_i = \sum_{i=1}^r v_i + \text{tx_max_fee},$$

where $\text{tx_max_fee} = \text{gas_total} \cdot \text{gas_price}$. When creating the transaction to transfer the funds, the sender will have to nullify the set of old notes being spent and mint a new set of notes $\{\text{N}_1^{\text{new}}, \dots, \text{N}_r^{\text{new}}\}$ with the corresponding values v_i , and assigned to the corresponding receivers. The most common case is $r = 2$, where a sender generates a note N_1^{new} with value v for a receiver, and a second note N_2^{new} for themselves with value $\text{change} = v - \sum_{i=1}^r w_i - \text{tx_max_fee}$. We detail the steps the sender should follow hereunder.

Mint new notes. Set the parameters of each new note $\text{N} \in \{\text{N}_i^{\text{new}}\}_{i=1}^r$ as follows.

1. Set the type of transaction.

³ The pair (npk, R) is sometimes referred to as a *stealth address* [34].

- If the transaction is transparent, set **type** = 0.
- If the transaction is obfuscated, set **type** = 1.
- 2. Set **v** to the amount of money being transferred to the new note **N**.
- 3. Set a blinding factor for the commitment and a nonce for the encryption.
 - If **type** = 0, set $s = 0$ and **nonce** = 0.
 - If **type** = 1, set $s \leftarrow \mathbb{F}_p$ and **nonce** $\leftarrow \mathbb{F}_p$.
- 4. Compute the value commitment **com** = $\text{Com}_{\text{ck}}(v; s)$ following Section 3.5.
- 5. Encrypt the opening of **com**.
 - If **type** = 0, then set **enc** = **v**.
 - If **type** = 1, then set $r \leftarrow \mathbb{F}_p$ and **enc** = $\text{Enc}_{k_{\text{DH}}}(\mathbf{v} || s; \text{nonce})$, where k_{DH} is the symmetric Diffie–Hellman key described in Section 4.1, and $\text{Enc}()$ is the encryption function from Section 3.4 and [22], that uses k_{DH} coordinates and the **nonce** for the initialization vector.
- 6. Compute the note public key **npk** of the receiver as described in Section 4.1.
- 7. Set the new note to

$$\mathbf{N} = \{\text{type}, \text{com}, \text{nonce}, \text{enc}, \text{npk}, R\}.$$

Nullify old notes. For each note $\mathbf{N} \in \{\mathbf{N}_i^{\text{old}}\}_{i=1}^s$, with its respective position **pos**, being spent, the sender must do the following computations.

1. Compute $\text{nullifier} = H^{\text{Poseidon}}(\text{npk}' || \text{pos})$, where npk' is the nullification key of **N** as described in Section 4.1, and **pos** is the note’s position in the Merkle tree of notes.
2. If **type** = 1, then decrypt the encrypted information **enc** using the symmetric Diffie–Hellman key described in Section 4.1 to get an opening $(\mathbf{w}; t)$ for **com**.
3. Retrieve the old note private key **nsk** as described in Section 4.1.
4. Compute the hash of the transaction skeleton, $\text{tx}_{\text{hash}} = H^{\text{BLAKE2b}}(\text{tx}_{\text{skeleton}})$.⁴
5. Use **nsk** to generate a signature **sig** on tx_{hash} by running the double-key Schnorr signature scheme described in Section 3.6.

Prove that the transaction is performed correctly. Compute a zero-knowledge proof using the circuit depicted in Figure 1 to prove that the following conditions hold.

1. *Membership:* the sender must prove that every $\mathbf{N} \in \{\mathbf{N}_i^{\text{old}}\}_{i=1}^s$ is included in the Merkle tree of notes. To do so, the sender takes the parameters $(\text{type}, \mathbf{w}, t, \text{npk})$ associated with **N** and uses them as inputs of the circuit. Inside the circuit, the box **commit()** computes the commitment **com** using **w** and **t**, and together with the rest of elements, computes hash $h = H(\text{type}, \text{com}, \text{npk})$ of the note. The sender also provides a Merkle proof to show that **h** is in a leaf of the Merkle tree. Finally, the circuit takes **merkle_proof**,

⁴ Note that we hash the skeleton, and not the whole payload, which does not include the positions of the newly minted notes nor the proof. This is because these values are not yet determined at this point. Also, this step does not change for each note, so it can be done just once.

`root`, `pos` and `h`, and verifies the Merkle proof in `verify_merkle_proof()`, i.e. it checks that it leads to the provided `root`. Observe that all these inputs are private and hence, the proof will not reveal which note is being spent, only that it belongs to the Merkle tree of notes.

2. *Ownership*: the sender must prove that they hold the note secret key `nsk` of every note $N \in \{N_i^{\text{old}}\}_{i=1}^s$. Instead of including their private key as an input to the circuit and computing `npk` inside, the sender proves (using the `verify_signature()` box inside the circuit) that they can sign a message with that key. In this case, they use the double-key Schnorr signature scheme to sign the hash of the transaction. Note that the signature can be kept as a private input of the circuit, since a correct execution of the circuit is proving that the sender holds a note's secret key associated with a note's public key that belongs to the Merkle tree of notes. Not only that, since the transaction's hash `tx_hash` is already a public input for the circuit, it binds the details of the transaction with that particular signature.
3. *Nullification*: the sender must prove that `nullifier` = $H(\text{npk}' || \text{pos})$. Note that the sender provides the nullification key `npk'` = `nskG'` as an input to the circuit and not the note secret key `nsk`. As we just explained, the double-key Schnorr signature guarantees that `npk'` is indeed `nskG'`. The result of the right `hash()` box is the nullifier, which is a public output circuit that is later included as part of the transaction.
4. *Minting*: the sender inputs the value `v` being transferred and the randomness `s`, and the circuit computes the corresponding value commitment. In this case, `com` is the public output that comes from the right `commit()` box.
5. *Balance integrity*: the `verify_balance()` box checks that

$$\sum_{i=1}^r w_i - \sum_{i=1}^s v_i - \text{tx_max_fee} = 0, \quad (1)$$

where `gas` is the maximum amount of gas that the sender is willing to pay for the transaction. Furthermore, since each input and output is represented by an element of \mathbb{F}_p , we need to prevent overflows that could lead to generating money out of thin air. For this reason, we include a range check on each value:

$$\begin{aligned} 0 \leq v_i \leq 2^{64} & \quad \forall i = 1, \dots, r, \\ 0 \leq w_i \leq 2^{64} & \quad \forall i = 1, \dots, s. \end{aligned}$$

On the one side, observe that the public inputs of the circuit from Fig. 1 reveal nothing about the details of the transaction. On the other side, observe that the most important piece of information that allows a user to spend a note is `nsk`, which is obtained using both parts of this secret key `sk`, but neither key is included as an input to the circuit. Instead, Phoenix uses the double-key Schnorr signature as proof that the user holds `nsk`. As a consequence, users could delegate the generation of the zero-knowledge proof to a partially trusted third party, that is, a *proof helper*. This delegation would require the user to entrust the old and new values of the notes to the proof helper, but not their secret

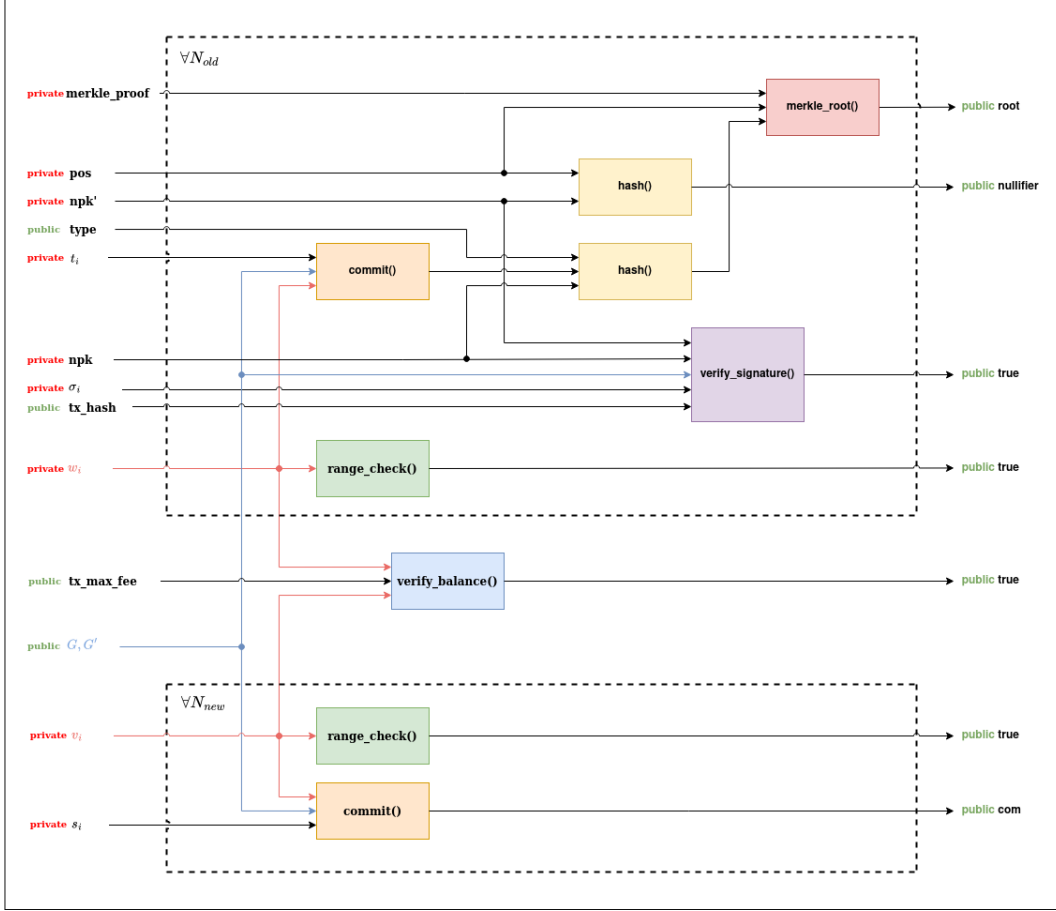


Fig. 1. Arithmetic circuit for an obfuscated Dusk transfer.

key. Benchmarks from Section 5.2 suggest that this disclosure is a reasonable trade-off with respect to the gained efficiency.

The remaining checks not verified inside the circuit are performed by the network. For instance, checking that the nullifier included in the transaction matches the output of the circuit, or that the note has the right typeset. Still, there are two computations the network cannot verify: checking that npk and enc are calculated correctly. If not, in both cases the receiver could not spend the note they own, either because they did not identify the transaction was sent to them (and can not compute nsk), or because they cannot decrypt the opening of the value commitment. Phoenix is a design in which senders can always reveal the details of their transactions at any point in time. Note that they only need to reveal the value being transferred, the randomness used in the commitment, and the public key of the receiver, but no compromised private information. Hence,

in these scenarios, the sender could always prove that they computed npk and enc correctly, either by disclosing the information or producing a zero-knowledge proof.

5 Analysis

We analyze different aspects of the Phoenix protocol described in Section 4.

5.1 Security discussion

Unforgeability. To prevent users from spending notes not in the Merkle tree (i.e. notes that do not come from an UTXO), senders are required to produce a proof that the notes being spent indeed belong to the tree. This is achieved through the Merkle proof that is part of the transaction circuit, and works as specified in Section 3.1 (`verify_merkle_proof()` in Figure 1).

Malicious helpers. Note that Phoenix allows for delegation of the work of listening to the network for transactions addressed at one, by sharing the view key $\text{vk} = (a, B)$ with a network-listening helper, as explained in Section 4.1. Even with that information, it is not possible to compute the note’s secret key nsk , so helpers cannot spend notes not belonging to the user being helped.

Additionally, proof computation can be delegated to a proof helper too (Section 4.2), by sharing the circuit inputs. Since nsk is not shared and cannot be computed from npk, npk' , this helper cannot spend the notes either.

Unlinkability. In the Phoenix transaction model, it is not possible to trace funds. That is, not even the sender of a note can tell when the receiver is spending the note. This is because each note uses a unique note public key. Furthermore, any revealing information is either encrypted (Section 3.4), committed in a hiding commitment (Section 3.5) or is a secret input of the transaction circuit, and thus is protected by the zero-knowledge property of the proof system (Section 3.2).

Transaction non-malleability. Phoenix prevents malicious attackers from modifying others’ transactions before they are added to the ledger. This is achieved by including the hash of the transaction as a public input to the circuit, and signing it (`verify_signature()` in Figure 1). We argue in two steps that this is enough to prevent malleability. First, an attacker cannot produce a proof of a false statement, due to the soundness of the proof system (Section 3.2). Therefore, the circuit is satisfied, which implies that the purported signature passes the verification. But then, a forged signature cannot pass the verification, due to the unforgeability of the signature scheme (Section 3.6).

Balance integrity. This property requires that no user can own more money than what he minted or received via payments from others. This is done by proving that Equation 1 holds (`verify_balance()` in Figure 1). Such proof cannot be forged due to the soundness of the proof system (Section 3.2). Note that this is done within the circuit so that the values themselves are not leaked.

Double-spending prevention. To prevent double spending of a note in the Merkle tree, sent to a public key npk with secret key nsk and stored at position pos , a transaction includes a special value called *nullifier*, computed as

$$\text{nullifier} = H(\text{nsk}G' \parallel \text{pos}).$$

For a note to be spent, the nullifier must have not been used before. Note that this value is deterministic on nsk , which allows the note to be spent, and pos . Hence, if anyone tries to spend the note again, they would get the same value and the network would detect that they are trying to double spend the note. To ensure that the nullifier is really the result of this operation, the sender is required to prove correctness of this computation (right `hash()` in Figure 1).

Here G' is used instead of G because $\text{nsk}G = \text{npk}$ can be computed by a proof helper. This helper would realize that the public nullifier corresponds to a certain npk . Thus, they would know when a certain note is being spent, breaking unlinkability.

Anonymity and privacy. The transaction does not contain any information about the sender, receiver or value of the transfer in the clear. The notes spent by the prover are kept as secret inputs of the proof, so leakage is prevented by the zero-knowledge property of the proof system (Section 3.2). A unique note public key $\text{npk} = H(rA)G + B$, which depends on the public key $\text{pk} = (A, B)$ of the receiver. However, due to the hash there is no way for an external attacker to link npk to pk . Finally, the value of the transfer appears only within the zero-knowledge proof, and in the commitment and encryption, which hide the information within (Sections 3.5 and 3.4, respectively).

Note that a proof helper does learn information like the values of the notes being spent and created.

5.2 Implementation and benchmarks

Dusk implemented [9] the Phoenix transaction model using PlonK as its zero-knowledge proof system. In particular, they use their own crate of PlonK [10]. Such libraries are coded in Rust, and are used by the different nodes of the network to generate the proof that comes out of the circuit depicted in Figure 1. Those proofs are later sent into a transaction, which calls the *transfer smart contract*. Such smart contract is executed and validated by the network validators, which will run an instance of Rusk [11], the architecture that builds the node workflows. This architecture includes a virtual machine, the Rusk VM, and one of its tasks is to execute the smart contracts, coded in Rust as well but compiled into WebAssembly. As such, the proofs computed by the users are verified by the transfer contract in the Rusk VM.

Depending on the number of notes a user wants to nullify, the implementation includes four different circuits (and thus, four different smart contracts), all of them resulting in two minted notes. In order to evaluate the benefits of delegating the generation of zero-knowledge proofs, we benchmark the proving time of

these four circuits on different hosts. As we can see in Table 5.2, computing the proof using a desktop CPU like the **Apple Silicon M1** can be done with good efficiency. On the other hand, a mobile CPU like the **Snapdragon 845** can raise the computing time up to 147 seconds. Besides, the amount of storage (approx. 3 GB) and RAM needed to execute the prover suggests that, as explained in Section 4, in many cases it is a better approach to delegate this task.

Circuit	Constraints	Proving time	
		Apple Silicon M1	Snapdragon 845
Nullify 1 note	31486	4.2083s	25.986s
Nullify 2 notes	61348	7.8706s	48.496s
Nullify 3 notes	91210	16.082s	98.804s
Nullify 4 notes	121072	18.431s	147.609s

Table 1. Benchmarks of Dusk’s PlonK implementation on each circuit available for the transfer contract. In all circuits, the amount of minted notes is 2. The table gathers the execution time of the prover on different hosts.

6 Conclusions

We reviewed the Phoenix transaction model implemented by Dusk. We showed that the privacy model is capable of issuing obfuscated transactions across Dusk’s blockchain by means of different cryptographic primitives. We confirmed the soundness of the protocol, analyzed its features, and benchmarked the performance of the protocol using their implementation, demonstrating its feasibility both in desktop and mobile hardware. Moreover, we explained the novel delegation model that allows users to delegate the heavy part of the transactions computations to partially trusted environments.

We have also seen how a different circuit for each amount of notes to be nullified and minted is required. This is a drawback as the storage required to store all the generated proving and verifying keys grows per each additional circuit we want to support. To solve this, an interesting future work would be to find a way to merge the notes to be nullified into one, before putting them into the circuit.

References

1. Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation) (Text with EEA relevance), available online: <http://data.europa.eu/eli/reg/2016/679/oj> (accessed on 1 June 2022)

2. AB-375 Privacy: personal information: businesses. An act to add Title 1.81.5 (commencing with Section 1798.100) to Part 4 of Division 3 of the Civil Code, relating to privacy. (2018), available online: https://leginfo.legislature.ca.gov/faces/billTextClient.xhtml?bill_id=201720180AB375 (accessed on 1 June 2022)
3. Bertoni, G., Daemen, J., Peeters, M., Van Assche, G.: Cryptographic sponges (2011)
4. Bonneau, J., Narayanan, A., Miller, A., Clark, J., Kroll, J.A., Felten, E.W.: Mixcoin: Anonymity for bitcoin with accountable mixes. In: International Conference on Financial Cryptography and Data Security. pp. 486–504. Springer (2014)
5. Bootle, J., Cerulli, A., Chaidos, P., Groth, J., Petit, C.: Efficient zero-knowledge arguments for arithmetic circuits in the discrete log setting. In: Annual International Conference on the Theory and Applications of Cryptographic Techniques. pp. 327–357. Springer (2016)
6. Bowe, S., Grigg, J.: Implementation of the BLS12-381 pairing-friendly elliptic curve construction Available online: https://github.com/zkcrypto/bls12_381 (accessed on 1 June 2022)
7. Bowe, S., Ogilvie-Wigley, E., , Grigg, J.: Implementation of the Jubjub elliptic curve group Available online: <https://github.com/zkcrypto/jubjub> (accessed on 1 June 2022)
8. Diffie, W., Hellman, M.: New directions in cryptography. IEEE Transactions on Information Theory **22**(6), 644–654 (1976)
9. Dusk Network: Phoenix core. GitHub, available online: <https://github.com/dusk-network/phoenix-core> (accessed on 1 June 2022)
10. Dusk Network: Plonk. GitHub, available online: <https://github.com/dusk-network/plonk> (accessed on 1 June 2022)
11. Dusk Network: Rusk. GitHub, available online: <https://github.com/dusk-network/rusk> (accessed on 1 June 2022)
12. Fiat, A., Shamir, A.: How to prove yourself: Practical solutions to identification and signature problems. In: Conference on the theory and application of cryptographic techniques. pp. 186–194. Springer (1986)
13. Fuchsbaauer, G., Kiltz, E., Loss, J.: The algebraic group model and its applications. In: Annual International Cryptology Conference. pp. 33–62. Springer (2018)
14. Gabizon, A., Williamson, Z.J., Ciobotaru, O.: PlonK: Permutations over lagrange-bases for ocumenical noninteractive arguments of knowledge. Cryptology ePrint Archive (2019)
15. Goodell, B., Noether, S., Blue, A.: Concise linkable ring signatures and forgery against adversarial keys. Cryptology ePrint Archive (2019)
16. Grassi, L., Khovratovich, D., Rechberger, C., Roy, A., Schofnegger, M.: Poseidon: A new hash function for zero-knowledge proof systems. In: 30th USENIX Security Symposium (USENIX Security 21). pp. 519–535 (2021)
17. Groth, J.: On the size of pairing-based non-interactive arguments. In: Annual international conference on the theory and applications of cryptographic techniques. pp. 305–326. Springer (2016)
18. Herrera-Joancomartí, J.: Research and challenges on bitcoin anonymity. In: Data privacy management, autonomous spontaneous security, and security assurance, pp. 3–16. Springer (2014)
19. Hopwood, D., Bowe, S., Hornby, T., Wilcox, N.: Zcash protocol specification. GitHub: San Francisco, CA, USA p. 1 (2016)
20. Kate, A., Zaverucha, G.M., Goldberg, I.: Constant-size commitments to polynomials and their applications. In: International conference on the theory and application of cryptology and information security. pp. 177–194. Springer (2010)

21. Katz, J., Lindell, Y.: Introduction to modern cryptography. CRC press (2020)
22. Khovratovich, D.: Encryption with Poseidon Available online: <https://dusk.network/uploads/Encryption-with-Poseidon.pdf> (accessed on 1 June 2022)
23. Maharramov, T.: Confidential security contract standard Available online: https://dusk.network/uploads/Confidential_Security_Contract_Standard_v2_0.pdf (accessed on 1 June 2022)
24. Maharramov, T., Khovratovich, D., Francioni, E.: The Dusk Network whitepaper Available online: https://dusk.network/uploads/The_Dusk_Network_Whitepaper_v3_0_0.pdf (accessed on 1 June 2022)
25. Merkle, R.C.: A digital signature based on a conventional encryption function. In: Conference on the theory and application of cryptographic techniques. pp. 369–378. Springer (1987)
26. Nakamoto, S.: Bitcoin: A peer-to-peer electronic cash system. Decentralized Business Review (2008)
27. Noether, S., Goodell, B.: Triptych: logarithmic-sized linkable ring signatures with applications. In: Data Privacy Management, Cryptocurrencies and Blockchain Technology, pp. 337–354. Springer (2020)
28. Pedersen, T.P.: Non-interactive and information-theoretic secure verifiable secret sharing. In: Annual international cryptology conference. pp. 129–140. Springer (1991)
29. Pointcheval, D., Stern, J.: Security proofs for signature schemes. In: International conference on the theory and applications of cryptographic techniques. pp. 387–398. Springer (1996)
30. Reid, F., Harrigan, M.: An analysis of anonymity in the bitcoin system. In: Security and privacy in social networks, pp. 197–223. Springer (2013)
31. Saarinen, M.J.O., Aumasson, J.P.: The BLAKE2 Cryptographic Hash and Message Authentication Code (MAC). RFC 7693 (Nov 2015), available online: <https://www.rfc-editor.org/info/rfc7693> (accessed on 1 June 2022)
32. Sasson, E.B., Chiesa, A., Garman, C., Green, M., Miers, I., Tromer, E., Virza, M.: Zerocash: Decentralized anonymous payments from bitcoin. In: 2014 IEEE symposium on security and privacy. pp. 459–474. IEEE (2014)
33. Schnorr, C.P.: Efficient identification and signatures for smart cards. In: Conference on the Theory and Application of Cryptology. pp. 239–252. Springer (1989)
34. Van Saberhagen, N.: Cryptonote v 2.0 (2013)
35. Victor Lopez: Rusk: Dusk genesis circuits. GitHub, available online: <https://github.com/dusk-network/rusk/blob/master/circuits/transfer/doc/dusk-genesis-circuits.pdf> (accessed on 12 January 2023)